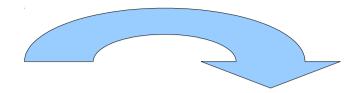
MLiB - Mandatory Project 3 Gene finding using HMMs

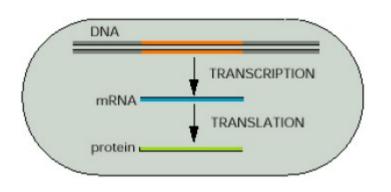
Viterbi decoding



>NC 002737.1 Streptococcus pyogenes M1 GAS TTGTTGATATTCTGTTTTTTCTTTTTTAGTTTTCCACATGAAAAATAGTTGAAAACAATA GCGGTGTCCCCTTAAAATGGCTTTTCCACAGGTTGTGGAGAACCCAAATTAACAGTGTTA ATTTATTTTCCACAGGTTGTGGAAAAACTAACTATTATCCATCGTTCTGTGGAAAACTAG AATAGTTTATGGTAGAATAGTTCTAGAATTATCCACAAGAAGGAACCTAGTATGACTGAA AATGAACAAATTTTTTGGAACAGGGTCTTGGAATTAGCTCAGAGTCAATTAAAACAGGCA ACTTATGAATTTTTTGTTCATGATGCCCGTCTATTAAAGGTCGATAAGCATATTGCAACT ATTTACTTAGATCAAATGAAAGAGCTCTTTTGGGAAAAAAATCTTAAAGATGTTATTCTT ACTGCTGGTTTTGAAGTTTATAACGCTCAAATTTCTGTTGACTATGTTTTCGAAGAAGAC CTAATGATTGAGCAAAATCAGACCAAAATCAACCAAAAACCTAAGCAGCAAGCCTTAAAT TCTTTGCCTACTGTTACTTCAGATTTAAACTCGAAATATAGTTTTGAAAACTTTATTCAA GGAGATGAAAATCGTTGGGCTGTTGCTGCTTCAATAGCAGTAGCTAATACTCCTGGAACT ACCTATAATCCTTTGTTTATTTGGGGTGGCCCTGGGCTTGGAAAAACCCATTTATTAAAT GCTATTGGTAATTCTGTACTATTAGAAAATCCAAATGCTCGAATTAAATATATCACAGCT GAAAACTTTATTAATGAGTTTGTTATCCATATTCGCCTTGATACCATGGATGAATTGAAA GAAAAATTTCGTAATTTAGATTTACTCCTTATTGATGATATCCAATCTTTAGCTAAAAAA ACGCTCTCTGGAACACAAGAAGAGTTCTTTAATACTTTTAATGCACTTCATAATAATAAC AAACAAATTGTCCTAACAAGCGACCGTACACCAGATCATCTCAATGATTTAGAAGATCGA TTAGTTACTCGTTTTAAATGGGGATTAACAGTCAATATCACACCTCCTGATTTTGAAACA CGAGTGGCTATTTTGACAAATAAAATTCAAGAATATAACTTTATTTTTCCTCAAGATACC ATTGAGTATTTGGCTGGTCAATTTGATTCTAATGTCAGAGATTTAGAAGGTGCCTTAAAA GATATTAGTCTGGTTGCTAATTTCAAACAAATTGACACGATTACTGTTGACATTGCTGCC GAAGCTATTCGCGCCAGAAAGCAAGATGGACCTAAAATGACAGTTATTCCCATCGAAGAA ATTCAAGCGCAAGTTGGAAAATTTTACGGTGTTACCGTCAAAGAAATTAAAGCTACTAAA CGAACACAAATATTGTTTTAGCAAGACAAGTAGCTATGTTTTTAGCACGTGAAATGACA GATAACAGTCTTCCTAAAATTGGAAAAGAATTTGGTGGCAGAGACCATTCAACAGTACTC CATGCCTATAATAAAATCAAAAACATGATCAGCCAGGACGAAAGCCTTAGGATCGAAATT GAAACCATAAAAAACAAAATTAAATAACATGTGGAAAAGAATATCTTTTATGAAATAGTT

>NC 002737.1 gene annotation Streptococcus pyogenes M1 GAS

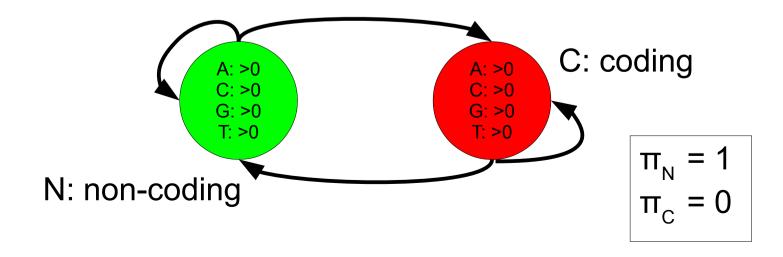
Design a HMM that models the syntax of genes

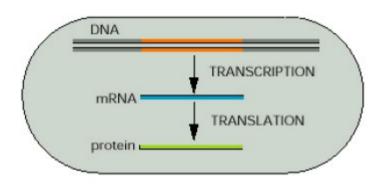


Biological facts

The gene is a substring of the DNA sequence of A,C,G,T's

X: acgatgcgctaatatgtccgatgacgtgagcataagcgacatgcag

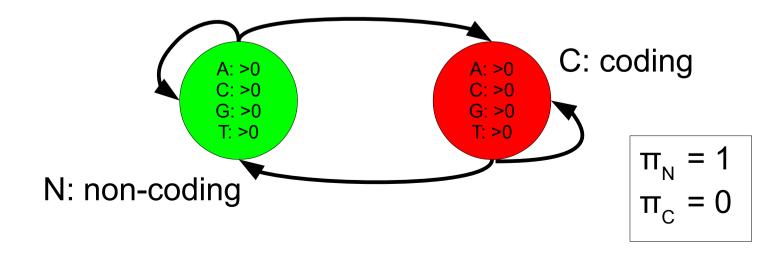


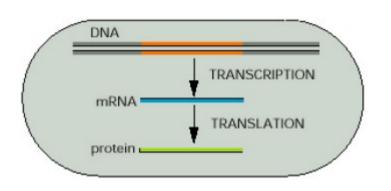


Biological facts

- The gene is a substring of the DNA sequence of A,C,G,T's
- The gene starts with a start-code atg

X: acgatgcgctaatatgtccgatgacgtgagcataagcgacatgcag





Biological facts

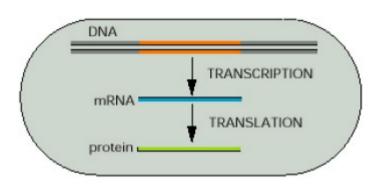
The gene is a substring of the DNA sequence of A,C,G,T's

C: coding

The gene starts with a start-code atg

X: acgatgcgctaatatgtccgatgacgtgagcataagcgacat $\pi_N = 1$ $\pi_C = 0$

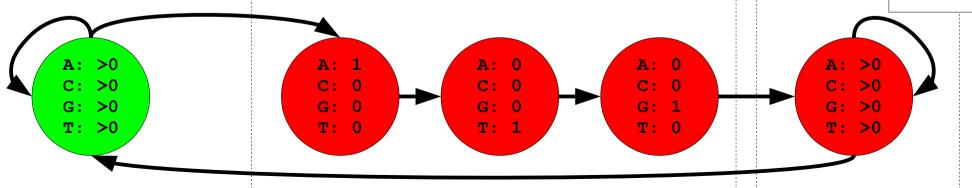
N: non-coding



Biological facts

- The gene is a substring of the DNA sequence of A,C,G,T's
- The gene starts with a start-code atg
- The gene ends with a stop-codon taa, tag or tga

X: acgatgcgctaatatgtccgatgacgtgagcataagcgacat $\pi_N = 1$ $\pi_C = 0$



N: non-coding

C: coding

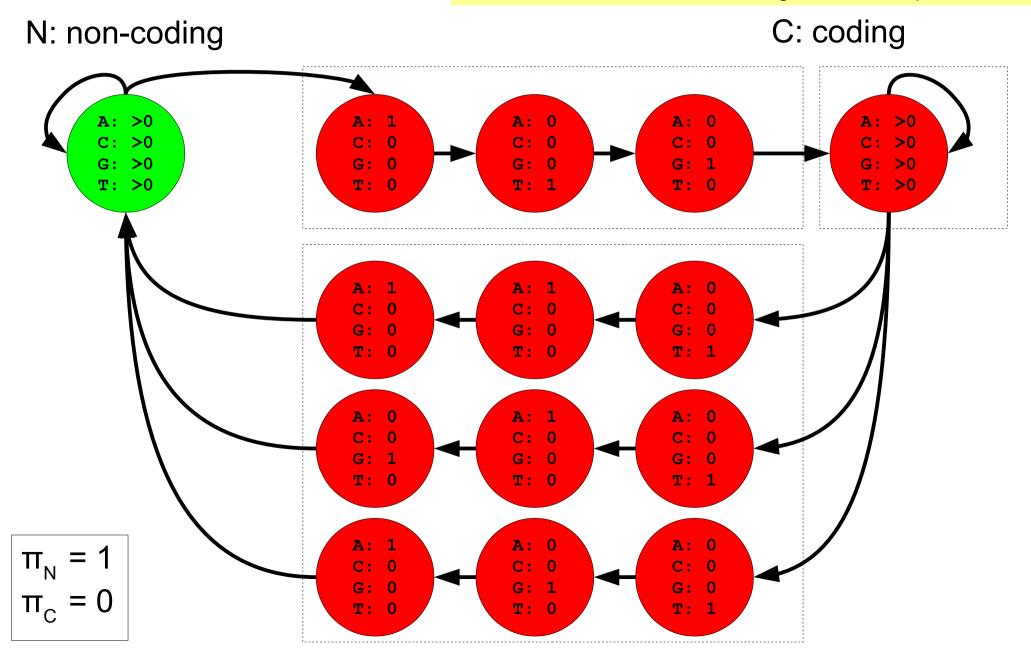
- The gene is a substring of the DNA sequence of A,C,G,T's
- The gene starts with a start-code atg

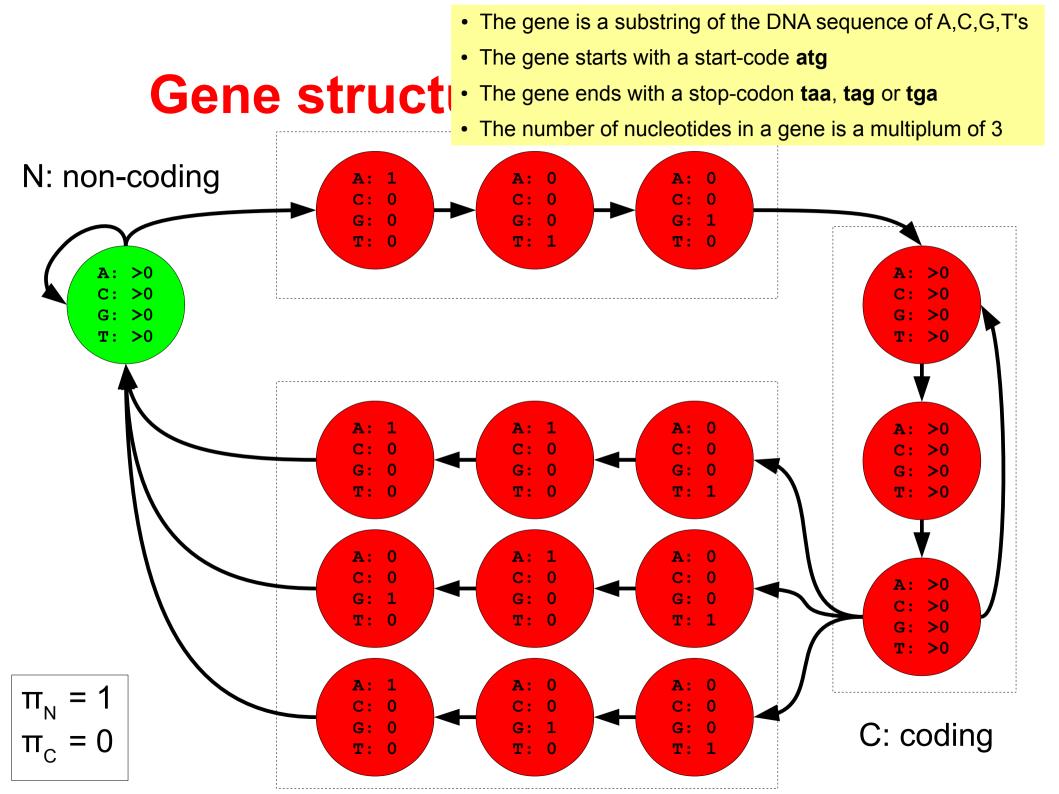
Gene struct • The gene ends with a stop-codon taa, tag or tga

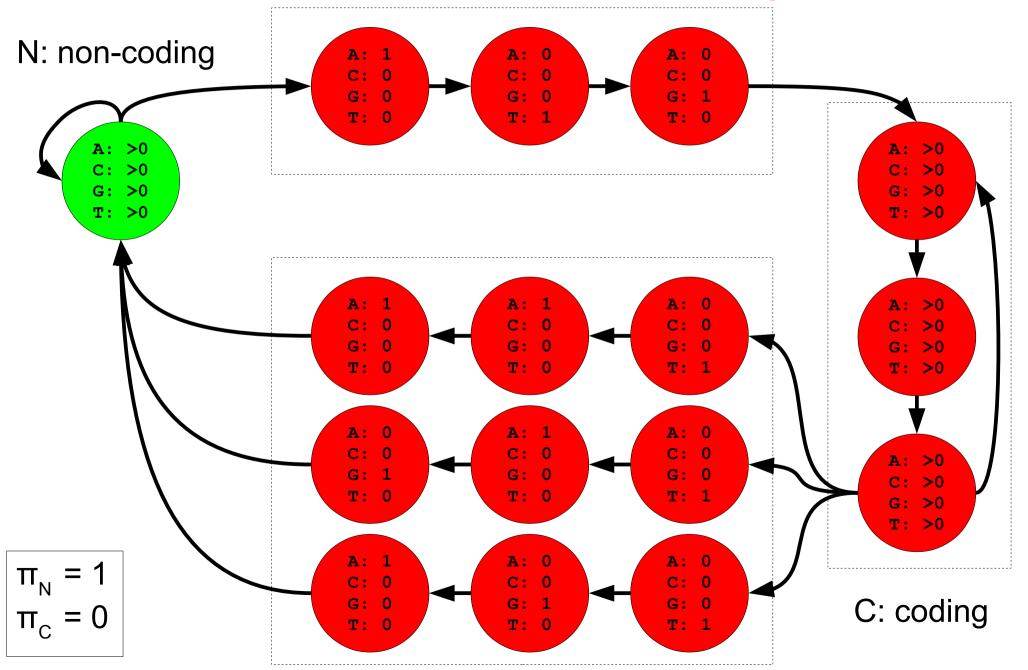
C: coding N: non-coding A: >0 A: >0 >0 C: 0 C: 0 C: 0 G: >0 G: >0 G: 0 G: 1 T: >0 T: >0 T: 0 T: 1 **T**: 0 A: 1 **A**: C: 0 **G**: 0 G: 0 **T**: 0 **T**: 0 T: 1 A: 0 **A**: 0 C: 0 C: 0 G: 1 G: 0 **T**: 0 T: 1 **T**: 0 **A:** 0 C: 0 C: 0 C: 0 G: 0 G: 0 **T**: 0 **T**: 0 T: 1

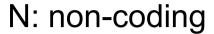
Gene struct • The gene ends with a stop-codon taa, tag or tga

- The gene is a substring of the DNA sequence of A,C,G,T's
- The gene starts with a start-code atg
- The number of nucleotides in a gene is a multiplum of 3









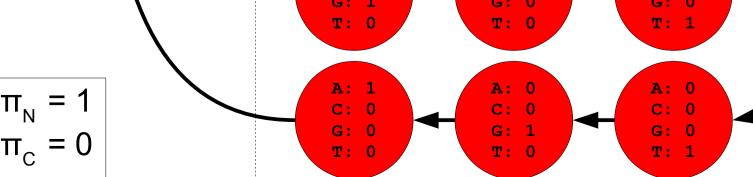
>0

G: >0

T: >0



- Select initial model structure (e.g. as done here)
- Select model parameters by training. Either "by counting" from examples of (**X**,**Z**)'s, i.e. genes with known structure, or by EM from examples of X, i.e. sequence which are known to contain a gene (as we will see on Monday)
- Given a new sequence X, predict its gene structure using the Viterbi algorithm for finding the most likely sequence of underlying latent states, i.e. its gene structure



C: coding

>0

C: >0

G: >0

T: >0

A: >0

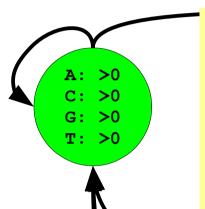
C: >0G: >0T: > 0

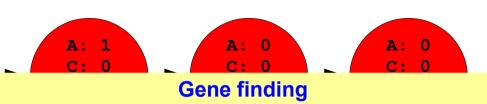
A: >0 C: >0

G: >0 T: > 0

Example – Gene finding

N: non-coding



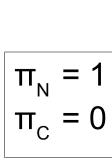


- Select initial model structure (e.g. as done here)
- Select model parameters by training. Either "by counting" from examples of (X,Z)'s, i.e. genes with known structure, or by EM from examples of X, i.e. sequence which are known to contain a gene (as we will see on Monday)

Even more biology

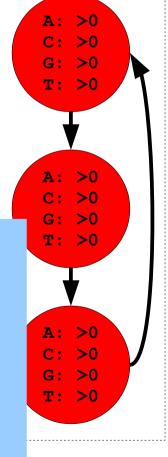
There can be genes in both directions (and over lapping)

- There are more possible start-codons atg, gtg, and ttg
- Internal codons cannot be start- or stop-codons
- And a lot more ...



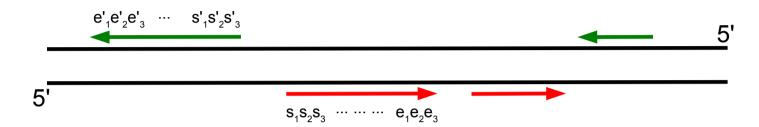
T: 0

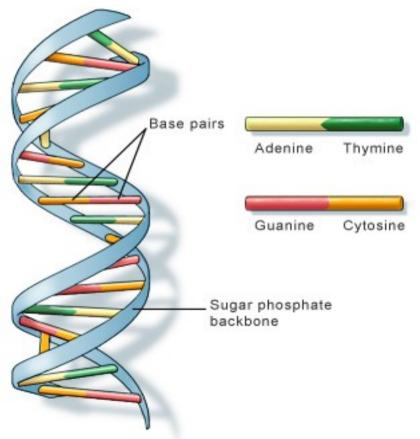
T: 1

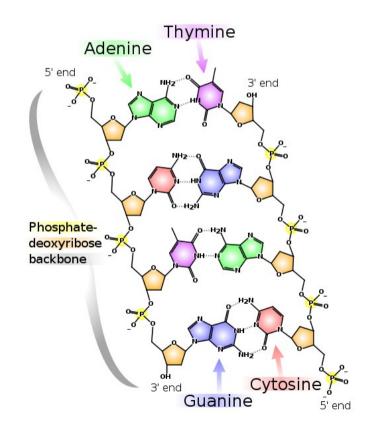


: coding

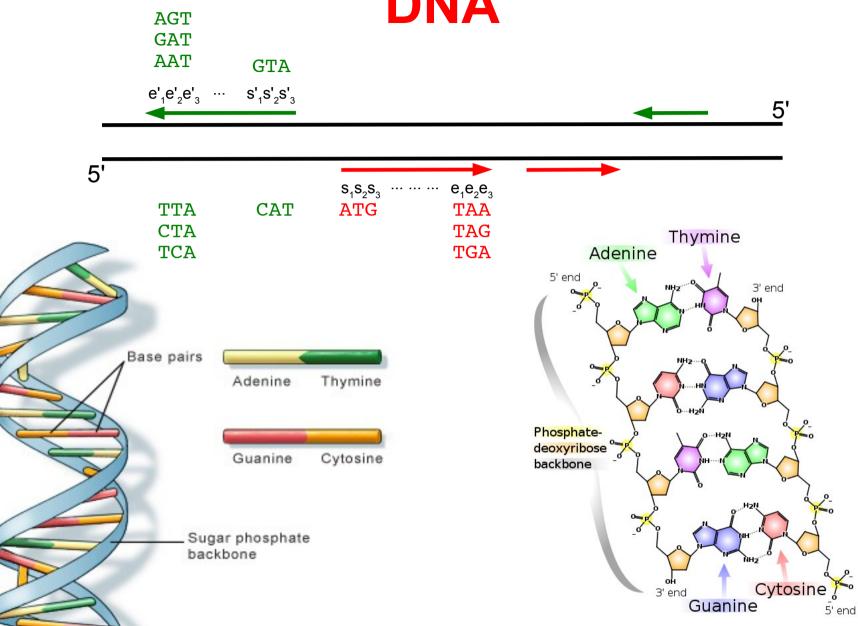
DNA







DNA



```
Length of genome1: 1852441 (1852441)
Length of genome2: 2211485 (2211485)
Length of genome3: 2499279 (2499279)
Length of genome4: 1796846 (1796846)
Length of genome5: 2685015 (2685015)
Length of genome6: 2127839 (2127839)
Length of genome7: 2742531 (2742531)
Length of genome8: 2046115 (2046115)
Length of genome9: 2388435 (2388435)
Length of genome10: 1570485 (1570485)
Length of genome11: 2096309 (2096309)
```

```
Start-codon in normal genes:
ATG [8423, 'NCCC']
ATC [3, 'NCCC']
ATA [1, 'RCCC']
GTG [713, 'NCCC']
ATT [3, 'NCCC']
CTG [2, 'NCCC']
GTT [1, 'NCCC']
CTC [1, 'NCCC']
TTA [1, 'NCCC']
TTG [1020, 'NCCC']
Stop-codon in normal genes:
TAG [1949, 'CCCN']
TGA [1531, 'CCCN']
TAA [6686, 'CCCN']
Reversed stop-codon in reversed genes:
TTA (reverse-complement: TAA) [6596, 'NRRR']
CTA (reverse-complement: TAG) [2014, 'NRRR']
TCA (reverse-complement: TGA) [1148, 'NRRR']
Reversed start-codon in reversed genes:
TAT (reverse-complement: ATA) [2, 'RRRN']
ATG (reverse-complement: CAT) [1, 'RRRN']
GAT (reverse-complement: ATC) [1, 'RRRN']
CAT (reverse-complement: ATG) [8077, 'RRRN']
AAT (reverse-complement: ATT) [4, 'RRRN']
TAC (reverse-complement: GTA) [1, 'RRRN']
CAC (reverse-complement: GTG) [715, 'RRRN']
CAA (reverse-complement: TTG) [953, 'RRRN']
CAG (reverse-complement: CTG) [4, 'RRRN']
```

```
Start-codon in normal genes:
ATG [8423, 'NCCC']
```

It might be an okay idea to ignore, i.e. not to model, the startand stop-codons which are very rare ...

This of course means that there might be genes that you Leng cannot predict ...

You also have to take care when "training by counting" because the data (X,Z) then contains annotations which are not possible in your model (you can e.g. skip start- and stopcodons which you do not model) ...

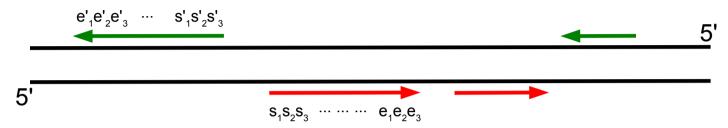
Length of genome11: 2096309 (2096309)

Leng Leng

Leng

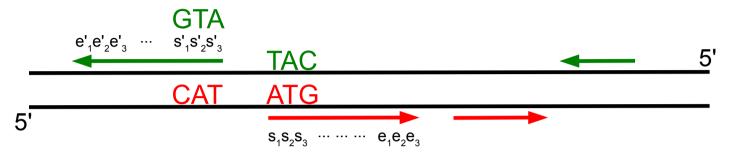
Leng

```
Reversed start-codon in reversed genes:
TAT (reverse-complement: ATA) [2, 'RRRN']
ATG (reverse-complement: CAT) [1, 'RRRN']
GAT (reverse-complement: ATC) [1, 'RRRN']
CAT (reverse-complement: ATG) [8077, 'RRRN']
AAT (reverse-complement: ATT) [4, 'RRRN']
TAC (reverse-complement: GTA) [1, 'RRRN']
CAC (reverse-complement: GTG) [715, 'RRRN']
CAA (reverse-complement: TTG) [953, 'RRRN']
CAG (reverse-complement: CTG) [4, 'RRRN']
```



The genes going from left-to-right (the C's):

- The first 3 symbols $s_1s_2s_3$ (the start-codon) is: **ATG**, ATC, ATA, **GTG**, ATT, CTG, GTT, CTC, TAA, **TTG**
- The last 3 symbols e₁e₂e₃ (the stop-codon) is: <u>TAG</u>, <u>TGA</u>, <u>TAA</u>
- The total length is a multiplum of 3.



The genes going from left-to-right (the C's):

- The first 3 symbols $s_1s_2s_3$ (the start-codon) is: **ATG**, ATC, ATA, **GTG**, ATT, CTG, GTT, CTC, TAA, **TTG**
- The last 3 symbols e₁e₂e₃ (the stop-codon) is: **TAG**, **TGA**, **TAA**
- The total length is a multiplum of 3.

The genes going from right-to-left (the R's):

The first 3 symbols s'₁s'₂s'₃ (the reversed start-codon) is: TAT, ATG, GAT, CAT, AAT, TAC, CAC, CAA, CAG

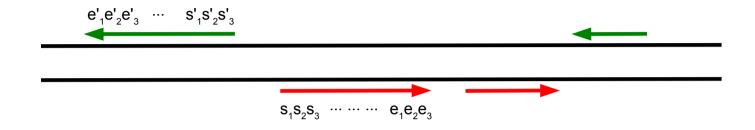
ATA, CAT, ATC, ATG, ATT, GTA, GTG, TTG, CTG

The last 3 symbols e'₁e'₂e'₃ (the reversed stop-codon) is: **TTA**, **CTA**, **TCA**

TAA, TAG, TGA

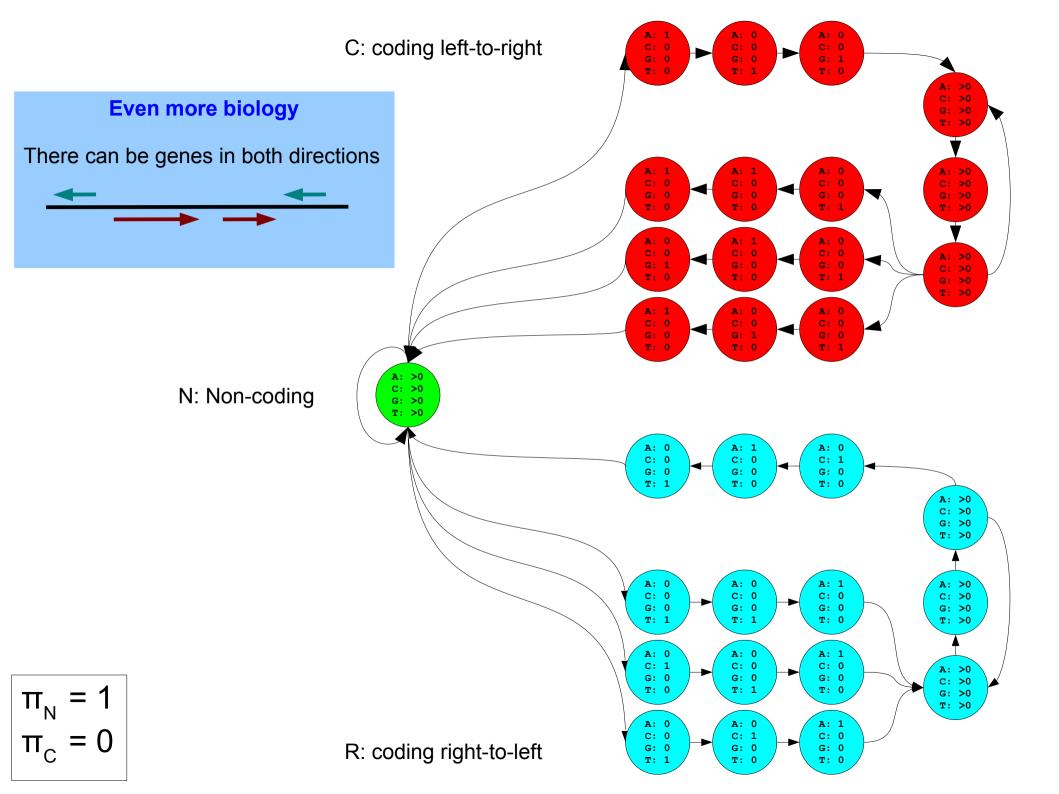
The total length is a multiplum of 3.

Predicting the full gene structure

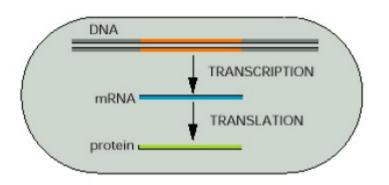


Approach for predicting the full gene structure (Cs and Rs)

- Predict C's and R's using two separate models, or the same model used for the genome (to predict C's) and its reverse complement (to predict R's).
- Afterwards we resolve conflicts, i.e. nucleotides predicted as both Cs and Rs.
- Or make a model which captures everything in one go.



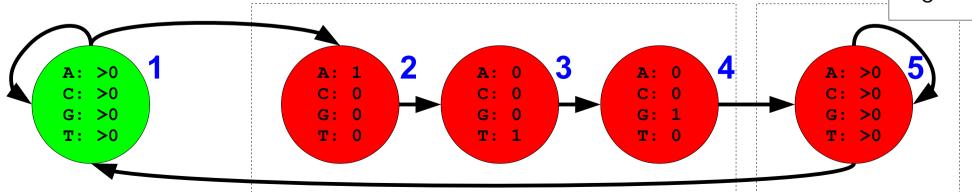
Problem: From annotation to Z



Problem: The string **Z=NNNCCC....** is not a prober sequence of states in the illustrated HMM, but is can easily be converted into one.

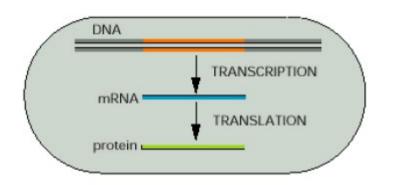
C: coding

X: acgatgcgctaatatgtccgatgacgtgagcataagcgacat $\pi_c = 0$



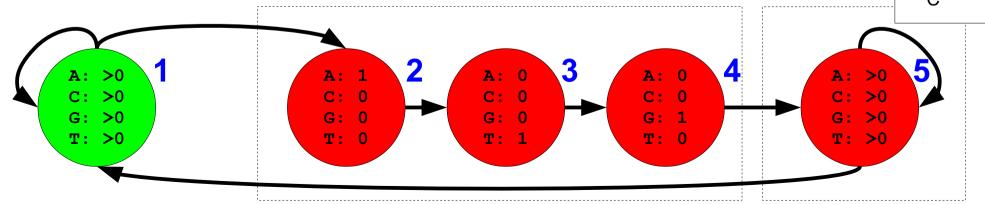
N: non-coding

Problem: From annotation to Z



Problem: The string **Z=NNNCCC....** is not a prober sequence of states in the illustrated HMM, but is can easily be converted into one.

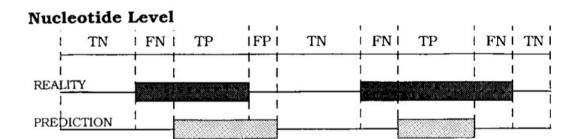
X: acgatgcgctaatatgtccgatgacgtgagcataagcgacat $\pi_c = 0$



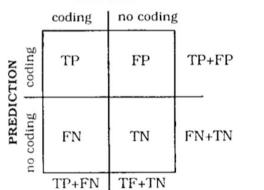
N: non-coding

C: coding

Evaluating performance







$$Sn = \frac{TP}{TP + FN}$$

Sensitivity

$$Sp = \frac{TP}{TP + FP}$$

Specificity

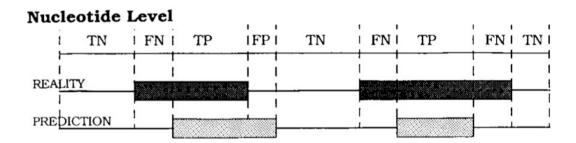
$$CC = \frac{(TP \times TN) - (FN \times FP)}{\sqrt{(TP + FN) \times (TN + FP) \times (TP + FP) \times (TN + FN)}}$$
 Correlation Coefficient

$$ACP = \frac{1}{4} \left[\frac{TP}{TP + FN} + \frac{TP}{TP + FP} + \frac{TN}{TN + FP} + \frac{TN}{TN + FN} \right]$$

$$AC = (ACP - 0.5) \times 2$$

 $AC = (ACP - 0.5) \times 2$ Approximate Correlation

Evaluating performance



Counting C's

```
def count_c(true, pred):
  total = tp = fp = tn = fn = 0
  for i in range(len(true)):
      if pred[i] == 'C' or pred[i] == 'c':
        total = total + 1
        if true[i] == 'C' or true[i] == 'c':
           tp = tp + 1
        else:
           fp = fp + 1
      if pred[i] == 'N' or pred[i] == 'n':
        if true[i] == 'N' or true[i] == 'n' or true[i] == 'R' or true[i] == 'r':
           tn = tn + 1
        else:
           fn = fn + 1
   return(total, tp, fp, tn, fn)
```

Counting C's

```
def count_c(true, pred):
  total = tp = fp = tn = fn = 0
  for i in range(len(true)):
      if pred[i] == 'C' or pred[i] == 'c':
        total = total + 1
        if true[i] == 'C' or true[i] == 'c':
           tp = tp + 1
        else:
           fp = fp + 1
      if pred[i] == 'N' or pred[i] == 'n':
        if true[i] == 'N' or true[i] == 'n'(or true[i] == 'R' or true[i] == 'r':
           tn = tn + 1
        else:
           fn = fn + 1
   return(total, tp, fp, tn, fn)
```

Counting R's

```
def count_r(true, pred):
  total = tp = fp = tn = fn = 0
  for i in range(len(true)):
      if pred[i] == 'R' or pred[i] == 'r':
        total = total + 1
        if true[i] == 'R' or true[i] == 'r':
           tp = tp + 1
        else:
           fp = fp + 1
      if pred[i] == 'N' or pred[i] == 'n':
        if true[i] == 'N' or true[i] == 'n' or true[i] == 'C' or true[i] == 'c':
           tn = tn + 1
        else:
           fn = fn + 1
   return(total, tp, fp, tn, fn)
```

Counting R's

```
def count_r(true, pred):
  total = tp = fp = tn = fn = 0
  for i in range(len(true)):
      if pred[i] == 'R' or pred[i] == 'r':
        total = total + 1
        if true[i] == 'R' or true[i] == 'r':
           tp = tp + 1
        else:
           fp = fp + 1
      if pred[i] == 'N' or pred[i] == 'n':
        if true[i] == 'N' or true[i] == 'n'(or true[i] == 'C' or true[i] == 'c':
           tn = tn + 1
        else:
           fn = fn + 1
   return(total, tp, fp, tn, fn)
```

Counting C's and R's

```
def count cr(true, pred):
  total = tp = fp = tn = fn = 0
  for i in range(len(true)):
     if pred[i] == 'C' or pred[i] == 'c' or pred[i] == 'R' or pred[i] == 'r':
        total = total + 1
        if (pred[i] == 'C' or pred[i] == 'c') and (true[i] == 'C' or true[i] == 'c'):
           tp = tp + 1
        elif (pred[i] == 'R' or pred[i] == 'r') and (true[i] == 'R' or true[i] == 'r'):
           tp = tp + 1
        else:
           fp = fp + 1
     if pred[i] == 'N' or pred[i] == 'n':
        if true[i] == 'N' or true[i] == 'n':
           tn = tn + 1
        else:
           fn = fn + 1
  return(total, tp, fp, tn, fn)
```